


— — Matched Dr.GRPO    — — Direct inverse conditional entropy    — — E53 inverse entropy + verified replay  
**Current canonical experiment — E55 ModeBench per-rollout verified-anchor sentinel — 0.5B**



Each row uses its own matched Dr.GRPO and preregistered evaluation. Every x-axis shows the full frozen 0-50-pass horizon; metric ranges tighten to the values observed so far. E55 uses K=8 exact-answer metrics with a 50-pass budget. Gold-support-normalized coverage is deliberately omitted from this decision view and never enters the controller or terminal gate. Cancelled predecessor treatments are omitted.

Bank admission is validator-bound in every treatment: Countdown uses the executed exact AST, graph coloring the constraint-checked vector, and Python factors the externally executed return vector. The model cannot supply its own canonical key.

E53 matches passive-tracking Dr.GRPO against direct inverse conditional-content entropy and the same discovery bridge plus teacher-forced verified-exemplar replay. Both inverse coefficients use only their own model-entropy warmup/EMA and have no coefficient bound. Replay capacity 16 is a compute bound; full verified discovery counts continue beyond it. Solid lines show the fresh E55 seed 9010; dashed controls are the frozen E53 matched Dr.GRPO seed 9010.