

Current canonical experiment — E55 ModeBench per-rollout verified-anchor sentinel — 0.5B

Python factors — E55 (live frontier)
graph coloring — E55 (live frontier)
Dr.GRPO and preregistered evaluation. Every x-axis expands with the latest treatment checkpoint; panels in a row share that live frontier, while metric ranges tighten to the values observed so far. E55 uses K=8 exact-answer metrics with a 50-pass sentinel budget, but live-chart ingestion remains uncapped and will display later checkpoints. Gold-support-normalized coverage is deliberately omitted from this decision view and never enters the controller or terminal gate. Cancelled prediction is retained. Direct conditional-content entropy remains active before any verified mode exists; uniform verified likelihood anchors the first observed correct exemplar; and only banks with at least two modes feed the canonical-entropy inverse controller. The anchor counts as one pseudo-rollout in the 16-sample group (fixed measure 1/16). Both inverse coefficients are unbounded, and training never reads gold support or evaluation. Solid lines show the fresh E55 seed 9010; dashed contr

