


Current canonical experiment — E57 ModeBench verified-first split-controller sentinel — 0.5B

Python factors — E57 (live frontier) — E57 (live frontier) — E57 (live frontier)



Dr.GRPO and preregistered evaluation. Every x-axis expands with the latest treatment checkpoint; panels in a row share that live frontier, while metric ranges tighten to the values observed so far. E57 uses K=8 exact-answer metrics with a 50-pass sentinel budget, but live-chart ingestion remains uncapped and will display later checkpoints. Gold-support-normalized coverage is deliberately omitted from this decision view and never enters the controller or terminal gate. Cancelled predictions are not counted. Bank admission is validator-bound in every treatment: Countdown uses the executed exact AST, graph coloring the constraint-checked vector, and Python factors the externally executed return vector. The model cannot supply its own canonical key. All groups have zero policy gradient and preserve the pretrained proposal distribution. After discovery, success-only open-set semantic pressure explores verified outcomes, verified surprisal controls common mass, and verified-bank entropy controls relative known-mode balance. All three coefficients use only their own 64-observation warmup/EMA state and remain unbounded. The replay losses share one fixed 1/16 pseudo-rollout measure; training never reads gold support or evaluation. Solids