


Five-domain confirmation — Dr.GRPO, historical E58, same-plumbing E66 control, and separated-support E68 actuator; 3 seeds × 12 passes
latest paper checkpoint=12.00 passes (training telemetry=12.00); complete 3-seed arm/domain points=103; fixed checkpoint cells=89/174; eval cadence=pass; audited E68 singleton escapes=327; thick lines are 3-seed means, shaded bands are seed ranges

● matched Dr.GRPO ◆ E58 global verified replay ■ E66 same-plumbing actuator-off control ▲ E68 separated-support entropy-gated actuator — seed 43 - - seed 44 ····· seed 45



ModeBench rows show their 10 fixed paper anchors. Held-out MATH-500 is evaluated every quarter epoch and this row shows the 7 registered even-pass anchors; its primary gate remains frozen to those same 7 anchors. Every row's x-axis expands as those points land. Training uses no gold support, desired mode count, desired entropy, or evaluation feedback; E58's three self-warmup coefficients remain unprojected. MATH-500 is held out; its mode-balance cells are N/A rather than inferred from prose. E66 is E68's same-plumbing causal comparator. E68 escape counts are support-only admissions and never PPO rows.