

INTERIM RESULTS — x-Mode Dr.GRPO across model scales

Qwen-0.5B 50/50; Falcon-1B 50/50; Qwen-3B 17/50 cells have registered evaluations; descriptive snapshot, final estimates pending
 semantic MaxEnt (eta = 0.1) cells with registered evaluations: Qwen2.5-0.5B with replay 25/25; Qwen2.5-0.5B no replay 25/25; Qwen2.5-0.5B adaptive 19/25; Qwen2.5-0.5B with replay 9/25; Falcon3-1B with replay 25/25
 Falcon3-1B no replay 22/25; Falcon3-1B adaptive 25/25; Qwen2.5-3B with replay 5/5; Qwen2.5-3B adaptive 5/5

