

INTERIM RESULTS — x-Mode GRPO across model scales

Qwen-0.5B/PointMaze 60/60; Falcon-1B 58/60; Qwen-3B 16/50 cells have registered evaluations; descriptive snapshot, final estimates pending
 semantic MaxEnt (eta = 0.1) cells with registered evaluations: Qwen2.5-0.5B with replay 25/25; Qwen2.5-0.5B no replay 25/25; Qwen2.5-0.5B adaptive 12/25; Qwen2.5-0.5B with replay 7/25
 Falcon3-1B with replay 23/25; Falcon3-1B no replay 15/25; Qwen2.5-3B with replay 5/5

