


dual — E36 matched Dr.GRPO (neutral; 3-seed mean) — E36 answer-option MI (neutral; 3-seed mean) - - E37 matched Dr.GRPO (3-seed mean) — E37 semantic collision entropy (3-seed mean) - · E38 predictive semantic Shannon entropy (3-seed mean) - - - E41 separately-centered semantic Sh

Countdown

Graph coloring

E32 rerun, E36 answer-option MI, and E37/E38/E41/E43 semantic-diversity evaluation

E32 rerun, E36 answer-option MI, and E37/E38/E41/E43 semantic-diversity evaluation



Rows: E32 EMA rerun, E36 neutral-prompt MI evaluation, and fresh E37 comparators with the E38 predictive-Shannon and E41 separately-centered Shannon-advantage extensions plus E43 success-conditioned signed Shannon advantage. Heavy lines require all three seeds; light bands show K=8 draw uncertainty.

training passes over the prompt pool