

A Verify
outputs

$y_1 \rightarrow c_1$ seen mode

$y_2 \rightarrow c_2$ new mode

$y_3 \rightarrow \perp$ invalid

one validator gives
reward and identity

B Store
exemplars

$\{c_1\}$ at start

↓ + c_2

$\{c_1, c_2\}$ kept

one exemplar per key,
in a lasting bank

C Revisit
recurrently

B1

B2

← τ

B3

one prompt-local bank
per optimizer step

D Replay
uniformly

$e_1 \in c_1$
 $e_2 \in c_2$



$L_{\text{rep}} = \text{mean}(-s\theta)$

every banked mode
keeps a gradient