

REWARD-ONLY PATH

POLICY SAMPLES



BINARY REWARD

reward: 1 1 1 0

mode identity discarded



MODE-BLIND UPDATE

A in 2 of 3 correct rows

B in 1 of 3 correct rows



SUPPORT NARROWS



B is gone from the batch

Re:Max (ours)

POLICY SAMPLES



REWARD + MODE

one validator returns both

reward: 1 1 1 0



reward
mode

MAXRL

fresh search learns



REPLAY
maintains



Re:Max UPDATE

MaxRL adds mode C

Replay keeps modes A, B



SUPPORT MAINTAINED + EXPANDED

A + B kept, C added

