

Correct-output concentration: all initial-to-final comparisons · Level 1




- Nominal stream representatives; nominal 95% seed interval
- Disjoint A lower / B upper
- Disjoint A upper / B lower

Graph



Qwen2.5-0.5B

Falcon3-1B




Qwen2.5-3B






Countdown





Python



MathIR



PantryPlan



-120 -60 0 60 120
ΔC (percentage points)

-120 -60 0 60 120
ΔC (percentage points)

-120 -60 0 60 120
ΔC (percentage points)

Positive ΔC: more concentrated correct outputs. Negative ΔC: less concentrated.
Labels show defined seed count and the primary eligible-prompt coverage range across seeds.
Open circles: fewer than five defined seeds. Sensitivity points use their own eligible prompts.
Correctness can change on these selected prompts; the companion table reports all original K8 metric changes.