


Correct-output concentration: training changes and replay effects · Level 1


- Nominal stream representatives; nominal 95% seed interval
- △ Disjoint A lower / B upper
- ▽ Disjoint A upper / B lower

Graph


Qwen2.5-0.5B






Falcon3-1B





Qwen2.5-3B






Countdown






Python



MathIR



PantryPlan



-120 -60 0 60 120

ΔC (percentage points)

-120 -60 0 60 120

ΔC (percentage points)

-120 -60 0 60 120

ΔC (percentage points)

Positive ΔC: more concentrated correct outputs. Negative ΔC: less concentrated.
Labels show defined seed count and the primary eligible-prompt coverage range across seeds.
Open circles: fewer than five defined seeds. Sensitivity points use their own eligible prompts.
Correctness can change on these selected prompts; the companion table reports all original K8 metric changes.